

## Model Details:

SBERT = type of transformer model trained specifically to create sentence embeddings

**Model name:** all-MiniLM-L6-v2

This is one specific SBERT model in the family.

### Details of the name:

all → trained on a variety of data sources, so it works for general-purpose text.

MiniLM → a smaller, faster transformer architecture.

L6 → 6 transformer layers (compact, good speed).

v2 → version 2 of this training setup.

**Trade-off:** fast (~384-dim vectors) but still accurate for semantic similarity.

With SBERT, each bio becomes a **384-dimensional vector of floats** learned by the model.

1. **Tokenize the text** → subwords (WordPiece/BPE).
2. **Transformer layers** encode context for each token.
3. **Pooling** turns all token vectors into one sentence vector.
4. **Normalize** (we set `normalize_embeddings=True`) → scale the vector to unit length so cosine similarity is just a dot product.

## How it works:

Sample Data:

row index (i)	user_id	age	sex	orientation	location	bio_text (short)
0	100	26	f	straight	NYC	"i like <b>hiking</b> and <b>music</b> "
1	101	27	m	straight	NYC	"i love <b>trails</b> and <b>concerts</b> "
2	102	24	m	gay	SF	"i like <b>hiking</b> and <b>coding</b> "
3	103	29	f	bi	LA	"i enjoy <b>swimming</b> and <b>coding</b> "

After SBERT encoding (and normalization), suppose the model gives these **3-D embeddings**

- E[0] (user 100): **A** = [0.7071, 0.7071, 0.0000]
- E[1] (user 101): **B** = [0.6000, 0.8000, 0.0000]
- E[2] (user 102): **C** = [0.7071, 0.0000, 0.7071]
- E[3] (user 103): **D** = [0.0000, 0.0000, 1.0000]

(All four are length 1, i.e., already normalized.)

When the `top_k_semantic_neighbors(user_id=100, k=2)` is run:

We're asking: *"Who are the top-2 semantic neighbors of user 100?"*

1) Find row index i

We look up where user\_id == 100:

So, i = 0 (first row)

2) Compute similarities `sims = E @ E[i]`

Because embeddings are normalized, dot product = cosine similarity.

Compute dot products with A:

$$\begin{aligned} A \cdot A &= 0.7071 \cdot 0.7071 + 0.7071 \cdot 0.7071 + 0 \cdot 0 \\ &= 0.5000 + 0.5000 + 0 = 1.0000 \end{aligned}$$

$$\begin{aligned} B \cdot A &= 0.6000 \cdot 0.7071 + 0.8000 \cdot 0.7071 + 0 \cdot 0 \\ &= 0.4243 + 0.5657 + 0 = 0.9900 (\approx 0.9899) \end{aligned}$$

$$\begin{aligned} C \cdot A &= 0.7071 \cdot 0.7071 + 0 \cdot 0.7071 + 0.7071 \cdot 0 \\ &= 0.5000 + 0 + 0 = 0.5000 \end{aligned}$$

$$\begin{aligned} D \cdot A &= 0 \cdot 0.7071 + 0 \cdot 0.7071 + 1 \cdot 0 \\ &= 0 + 0 + 0 = 0.0000 \end{aligned}$$

So sims = [1.0000, 0.9900, 0.5000, 0.0000]

3) Exclude self

`sims[i] = -1`

Now sims = [-1.0000, 0.9900, 0.5000, 0.0000].

4) Grab the top-k indices (fast)

`idx = np.argpartition(-sims, k)[k]`

We negate sims so “largest similarity” becomes “smallest negative number.” For k=2, this returns the indices of the two highest sims. Top two indices are 1 (0.99) and 2 (0.50).

5) Sort those k by true similarity

`idx = idx[np.argsort(-sims[idx])]`

This sorts the selected indices in descending similarity. idx = [1, 2] (since 0.99 > 0.50).

6) Build the output table

user_id	age	sex	orientation	location	score (cosine)
101	27	m	straight	NYC	0.9900
102	24	m	gay	SF	0.5000

Top-2 semantic neighbors for user 100 based on bio\_text meaning alone.